

THE UNIVERSITY of New South Wales

1st SEMESTER EXAMINATIONS

June/July 1999

DATABASE SYSTEMS - COMP9311

Solutions

1. (1) False.
(2) False.
(3) True.
(4) True.
(5) False.
(6) True.
(7) False.
(8) False.
(9) True.
(10) True.

2. (a)
 interface Person {
 attribute string name;
 attribute string birth_date;
 relationship Person motherOf
 inverse Person::childrenOfFemale
 relationship Person fatherOf
 inverse Person::childrenOfMale
 relationship Set<Person> children
 inverse Person::parentsOf
 relationship Set<Person> childrenOfFemale
 inverse Person::motherOf
 relationship Set<Person> childrenOfMale
 inverse Person::fatherOf
 relationship Set<Person> parentsOf
 inverse Person::children
 }

Notice that there are six different relationships here. For example, the inverse of the relationship that connects a person to their (unique) mother is a relationship that

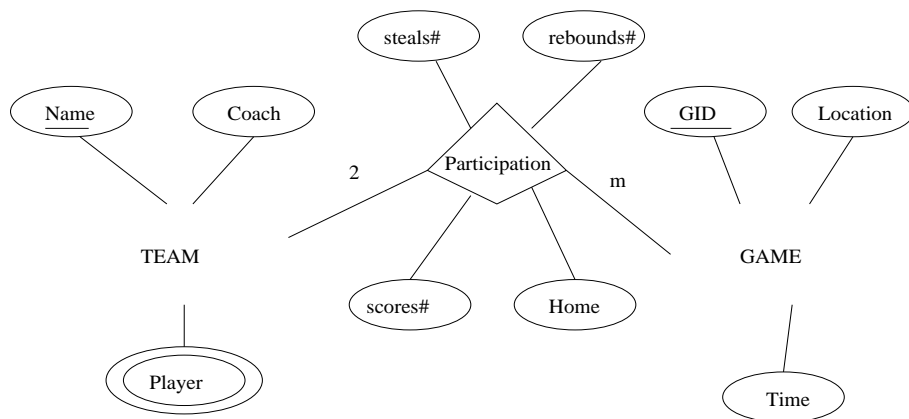


Figure 1

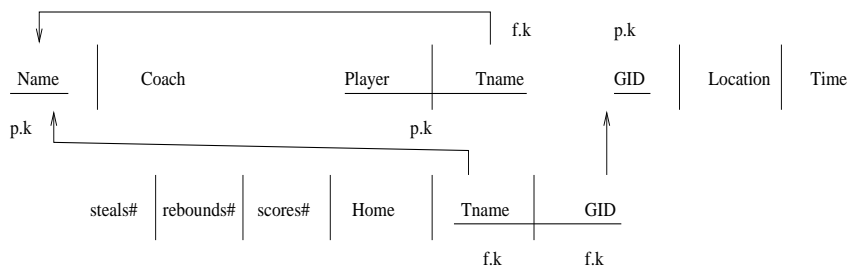


Figure 2

connects a mother (i.e., a female person) to the set of her children. That relationship, which we call `childrenOfFemale`, is different from the `children` relationship, which connects anyone – male or female – to their children.

(b.i) See Figure 1.

(b.ii) See Figure 2.

3. (a.i) Check all three constraints.

(a.ii) Referential constraints.

(a.iii) None.

(b) No. Because $(AC)^+ = \{A, B, C, D, H\}$ and $G \notin (AC)^+$.

(c) 1st NF. Because AB is the only key and C is partially dependent on A .

(d) $\{Y \rightarrow X, Z \rightarrow Y, Z \rightarrow W\}$.

4. (a.i) A primary index is an index on a set of fields that includes the primary key. It is guaranteed not to contain duplicates. A secondary index is an index that is not a primary index. It can have duplicates.

(a.ii) A dense index has at least one data entry search key value that appears in a record in the indexed file. A sparse index contains an entry for each page of records in a




Figure 4

data file. It *must* be clustered. Thus, we can have at most one sparse index on a data file. A sparse index is typically much smaller than a dense index.

- (a.iii) A clustered index is one in which the ordering of data entries is the same as the ordering of data records. We can have at most one clustered index on a data file. An unclustered index is an index that is not clustered. We can have several unclustered indexes on a data file.
 - (b) No, it is not a good idea, because it does not provide useful indexing information for retrieval and is useless.
 - (c) No. In this case, the index is unclustered, each qualifying data entries could contain an rid that points to a distinct data page, leading to as many data page I/O as the number of data entries that match the range query. In this case, if the number of data entries is larger than the number of data blocks then using index is worse than file scan.
5. (a) See my lecture notes for such an example.
- (b.i) A naive way to do this is to construct a serial schedule using the two-phase locking protocol.
 - (b.ii) No, it is impossible to have a non-serial but serializable schedule. Consider that the last operation in T_1 is WRITE (B), and the last operation in T_2 is WRITE (A). Meanwhile, T_1 starts with READ (A), and T_2 starts with READ (B). Therefore, in any serializable schedule of T_1 and T_2 , either READ (A) in T_1 should be after WRITE (B) in T_2 , or READ (B) in T_2 should be after WRITE (B) in T_1 .
 - (b.iii) Yes. (Students must construct a schedule with deadlock. It is very straightforward to construct such a schedule, so I omit it here.)
 - (c) The schedule is not serializable, because there is a cycle in its schedule graph. See Figure 3.